GANJA NILESH

Amalapuram, Andhra Pradesh ● +91-9381541130 ● nilesh30062004@gmail.com
LinkedIn ● GitHub ● LeetCode ● Portfolio

SUMMARY

AI Engineer with hands-on experience building multi-agent architectures, multimodal agentic pipelines, RAG frameworks, and LLM evaluation frameworks. Shipped a **24-agent multi-agent memory system** running at **90% production accuracy** and led cost optimization cutting LLM spend by **92%**.

EXPERIENCE

Kookar AI | LOR **08 2025 – 05 2026**
AI Product Analyst Intern *Gurugram, India*

- Engineered production LLM workflows for the Preferences & Memory pod across **3 LLM workflows** (Onboarding, Chat, Cook) and **24 production agents**, bridging product discovery with system architecture.
- Designed and shipped a multi-agent memory system spanning **6 memory types**: customer-side (meal, grocery, operations, persona) and cook-side (cook-operations, cook-others). Captures food, grocery, persona, and operational preferences continuously from onboarding calls and WhatsApp chats (text, audio notes, images, links).
- Built and optimized a Customer Approval Agent orchestrating **4 approval workflows** meal-plan approval, grocery-cart approval, order placement, and wallet recharge nudges with engagement-based triggers; improved approval and order-completion rates by **2**× and reduced approval time by **67%**.
- Executed household onboarding workflows for **15+ households**, coordinating activation activities that supported **10+ customers** in progressing from D7 to D14 retention.
- Established LangSmith LLM observability and Root Cause Analysis (RCA) tracing pipelines to resolve production failures, driving SOP refinement loops to hit **90% production accuracy**.
- Led memory-pod cost optimization by migrating 24 agents from GPT, Claude Opus & Sonnet to lite LLMs (GLM 4.7, DeepSeek v3.2 / v4 Pro, GLM 5, Kimi), cutting **LLM spend by 92%** at **90% accuracy**.

PROJECTS

GEO Auditor | *Python, FastAPI, BeautifulSoup4, httpx, NLP Heuristics* **08 2026 – Present**

- Built an automated Generative Engine Optimization audit API that evaluates website AI-readiness across **4 deterministic checks** using targeted crawling and content analysis.
- Implemented HTTP crawling with ‘httpx‘, ‘BeautifulSoup4‘ HTML/JSON-LD parsing, ‘robots.txt‘/‘llms.txt‘ analysis, and NLP heuristics for extractability signals.
- Developed a weighted **0–100** AI-readiness score with evidence-backed findings and prioritized remediation plans.

Multimodal AI Agent | *Python, FastAPI, Groq Whisper, PyMuPDF, Tesseract, Pydantic, Docker* **07 2026 – 08 2026**

- Engineered a custom multi-stage agentic orchestrator (Clarifier, Planner, Executor, Composer) in pure Python without LangChain, using Pydantic validation and retry logic for strict structural output bounds.
- Built an ingestion pipeline combining PyMuPDF with Tesseract OCR fallback for scanned PDFs, Tesseract for image text extraction, and Groq Whisper (whisper-large-v3-turbo) for audio transcription.
- Designed a **3-panel UI** displaying extracted files, markdown answers, and execution plan traces; deployed live.

AskMyDocs: Chat-with-PDF RAG App | *Python, RAG, FAISS, BM25, Cross-encoder, RAGAS, LangSmith* **05 2026 – 06 2026**

- **User-facing web app** that lets users upload PDFs and ask natural-language questions with cited page references. Engineered production RAG pipeline: hybrid retrieval (BM25 + dense + RRF), cross-encoder reranking, agentic 8B-router + Gemini 2.5 Flash answerer, token streaming. LangSmith-instrumented (dev/prod project split) and RAGAS-evaluated; **1.00 context recall**, **0.67 faithfulness**, **83% page-hit** on **10-Q golden set**; deployed live.

Font Recognition Using CNNs | *Python, Deep Learning, CNN, LeNet* **01 2025 – 05 2025**

- CNN (LeNet) for OCR font recognition; reduced overfitting by **20%** via Early Stopping + LR scheduling. **93.3%** font-style classification, **0.03px MAE** on font-size regression, **95%** combined accuracy.

TECHNICAL SKILLS

AI / LLM / GenAI : LLMs, GenAI, RAG, Agents, Multimodal AI, Prompt Engineering, LLM Evaluation, LangGraph, LangChain, FAISS, BM25, RAGAS, LangSmith, LLM Observability, Model Context Protocol (MCP)
ML / DL / Frameworks : Python, SQL, FastAPI, Docker, pytest, Pydantic, TensorFlow, Keras, scikit-learn, CNNs, Machine Learning, Deep Learning, Git
Product & Agile : PRDs, Agent Workflow Design, RCA, SOP Authoring, A/B Iteration, Stakeholder Communication
CS Fundamentals : Data Structures & Algorithms, Operating Systems, DBMS

EDUCATION

IIIT Sricity **08 2022 – 06 2026**
B.Tech in Computer Science & Engineering | ***CGPA: 8.13 / 10*** *Sri City, Andhra Pradesh, India*

ACHIEVEMENTS

- Scored **87%** in TCS NQT February 2025 (National Qualifier Test)
- Ranked **Top 10** out of **300 teams** in Thrive Digital National Healthcare Hackathon 2023